

Training-Free Generation of Protein Sequences from Small Family Alignments via Stochastic Attention

Jeffrey D. Varner^{a,1}

^aR.F. Smith School of Chemical and Biomolecular Engineering, Cornell University, Ithaca, NY 14850

This manuscript was compiled on March 17, 2026

1 Generating novel protein sequences that respect the statistical constraints of a family typically requires training deep generative models on
2 thousands to millions of examples. Yet most protein families are small: the median Pfam family has fewer than 100 members, a regime where
3 learned models overfit or collapse. We propose *stochastic attention* (SA), a training-free sampler that treats the modern Hopfield energy
4 over stored protein sequences as a Boltzmann distribution and draws samples via Langevin dynamics. The score function is given in closed
5 form by the residual of a single softmax attention operation, eliminating the need for a trained score network, external pretraining data, or
6 graphics processing unit (GPU) resources. Across eight Pfam families spanning 37 to 420 sequences and 23 to 262 residues, SA generates
7 sequences with low amino acid composition divergence, substantial novelty, and structural plausibility confirmed by ESMFold and AlphaFold2.
8 Head-to-head comparison with profile HMMs, EvoDiff, and the MSA Transformer shows that SA is the only method to simultaneously achieve
9 low composition divergence, genuine novelty, and sequence identity within the natural within-family range (30–70%); competing methods
10 either drift outside the family or produce near-copies of stored patterns. The critical inverse temperature is predicted from principal component
11 analysis (PCA) dimensionality alone with high accuracy, enabling fully automatic operation from a seed alignment. Independent validation by
12 ESM2-650M and cross-referencing against deep mutational scanning experiments confirm that over 90 percent of SA-generated substitutions
13 correspond to experimentally tolerated mutations.

protein sequence generation | stochastic attention | Hopfield networks | Langevin dynamics | training-free generative models

1 With protein structure prediction largely solved by AlphaFold2 (1) and RoseTTAFold (2), the frontier
2 has shifted toward *designing* novel protein sequences with desired properties (3, 4). Generative
3 approaches now span a broad landscape, from classical profile hidden Markov models (pHMMs) (5, 6) and
4 Potts models (7–9) to deep generative models including variational autoencoders (10, 11), autoregressive
5 language models (12, 13), alignment-processing transformers (14), MSA-conditioned diffusion (15), and
6 structure-conditioned methods such as ProteinMPNN (16) and RFdiffusion (17). Yet all sequence-level
7 approaches share a common requirement: large training corpora. Language models are pretrained on
8 millions of UniRef50 sequences; even MSA-conditioned methods require either extensive pretraining or large
9 input alignments. This reliance on scale creates a fundamental blind spot, because the median Pfam family
10 contains fewer than 100 full-length members (18) and thousands of families have fewer than 50. In this
11 scarce-data regime, deep models overfit and collapse in diversity, while pHMMs emit sequences that drift
12 far from the family’s functional envelope. What is needed is an approach that can extract the statistical
13 structure of a protein family from whatever data is available, even a few dozen sequences, without learning
14 parameters that could overfit.

15 A separate line of work provides exactly this capability. Ramsauer et al. (19) showed that a single softmax
16 attention operation is equivalent to one step of the retrieval dynamics on the modern Hopfield energy,
17 building on the dense associative memory framework of Krotov and Hopfield (20). This connection endows
18 any set of stored examples with a well-defined energy landscape: the examples become memory patterns,
19 and the Boltzmann distribution over this landscape assigns high probability to states that resemble the
20 stored set. Critically, the Hopfield energy provides the score function *analytically*, as the residual of a single
21 attention operation, so Langevin sampling can proceed without learning any parameters. We exploit this
22 insight to propose *stochastic attention* (SA), a training-free protein sequence generator that combines the
23 exact Hopfield score function with Langevin dynamics. Our prior work (21) introduced SA and validated it
24 on synthetic data, MNIST, and a single protein family.

25 In this study, we explored the extent to which stochastic attention can serve as a practical generative

model for real protein families. We systematically evaluated SA across eight diverse Pfam families spanning a ten-fold range of family sizes and a seven-fold range of sequence lengths, assessing compositional fidelity, sequence novelty, and structural plausibility using both ESMFold and AlphaFold2. We characterized how generation quality scales with family size, uncovering an empirical law that predicts the critical temperature from alignment statistics alone and eliminates the need for a temperature sweep. We tested robustness in the extreme scarce-data regime down to $K=20$ sequences, compared SA head-to-head against four established baselines (profile HMMs, EvoDiff, MSA Transformer, and a Potts model), and examined whether the generated sequences preserve the pairwise residue covariation that reflects structural contacts and functional couplings. Finally, we sought independent validation of biological plausibility through ESM2-650M pseudo-perplexity scoring and cross-referencing against published deep mutational scanning datasets.

Results

We ran stochastic attention on all eight Pfam families (Table 1) and observed that the method consistently generated novel protein sequences while preserving family-level amino acid statistics, across a ten-fold range of family sizes and a seven-fold range of sequence lengths (Fig. 1; SI Appendix, Table S1). We first confirmed that the sampler converged to the correct Boltzmann distribution by comparing the outputs of the unadjusted Langevin algorithm (ULA) and its Metropolis-adjusted variant (MALA): acceptance rates exceeded 99.6% in every family (SI Appendix, Table S7), confirming negligible discretization bias and justifying the use of the computationally cheaper ULA in all subsequent experiments. Generated sequences were not near-duplicates of stored patterns: zero sequences across all eight families exceeded 90% identity to any stored sequence (SI Appendix, Table S9), and the generated ensemble was insensitive to initialization strategy (SI Appendix, Table S8). In the generation regime (inverse temperature $\beta \approx 2\beta^*$, just above the entropy inflection point where pattern-specific basins first emerge), SA achieved Kullback–Leibler (KL) divergences below 0.06 in every family, with several below 0.02. Novelty ranged from 0.40 (Defensin_beta) to 0.65 (WW), indicating substantial departure from any single stored memory, and mean pairwise cosine distance was ≈ 1.0 in every family, confirming that the sampler produced distinct, well-separated sequences rather than collapsing to a single mode. Nearest sequence identity to a stored pattern was 51–66%, comparable to the typical within-family identity range, indicating that generated sequences remained within the family’s functional envelope while departing meaningfully from stored patterns. By contrast, bootstrap samples had zero novelty by construction and Gaussian perturbation produced negligible novelty (< 0.02), confirming that small additive noise is insufficient to escape the basin of the nearest stored pattern.

To understand why SA preserves these statistics, we examined the position-specific behavior of the sampler in the Kunitz domain (SI Appendix, Fig. S5). SA-generated sequences preserved all 11 positions with greater than 90% conservation in the seed alignment, including all six cysteines forming the three disulfide bonds essential to the Kunitz fold, while introducing 17–25 substitutions concentrated at variable positions. Per-position Shannon entropy correlated strongly between the SA-generated ensemble and the stored alignment ($r=0.968$), confirming that the method correctly identifies which positions tolerate variation and which do not. This behavior is a direct consequence of the Boltzmann sampling mechanism: at strongly conserved positions, all stored patterns agree and the softmax weights concentrate probability on the consensus residue; at variable positions, stored patterns diverge and the sampler draws from the local distribution of tolerated amino acids.

To characterize how generation quality degrades with decreasing family size, we subsampled the WW domain at $K \in \{20, 50, 100, 200, 400\}$ with five independent replicates per condition (SI Appendix, Fig. S2). Stochastic attention maintained low KL divergence across the entire range (0.019 ± 0.008 at $K=20$, improving to 0.010 ± 0.002 at $K=400$), while the PCA dimensionality grew from $d=17$ –18 components at $K=20$ to $d=182$ –183 at $K=400$ and the critical temperature tracked this growth ($\beta^*=2.7$ to 5.5). Despite the roughly ten-fold increase in effective dimensionality, SA generation quality remained stable, suggesting that the entropy inflection criterion successfully adapted the operating temperature to the geometry of the memory

74 matrix at each scale. Novelty increased with K (from 0.47 to 0.74), reflecting the richer landscape of a
75 larger memory set, while nearest sequence identity decreased correspondingly (from 0.71 to 0.56), indicating
76 that the sampler explored further from stored patterns when more patterns were available to collectively
77 shape the energy landscape. Replication across three additional families at $K=20$ (SH3, Kunitz, zf-C2H2)
78 confirmed that this robustness is not family-specific (SI Appendix, Table S2).

79 The scaling study revealed that β^* tracks PCA dimensionality; we next asked whether this relationship
80 could be predicted from alignment statistics alone, without running an explicit temperature sweep (Fig. 2).
81 We pooled 33 data points (the 8 Pfam families and 25 WW scaling replicates) and found that a simple
82 two-parameter model, $\beta^* \approx 1.57 + 0.28\sqrt{d}$ (bootstrap standard error [SE]: intercept ± 0.07 , slope ± 0.007),
83 explained 97% of the variance ($R^2=0.97$). The intercept of 1.57 recovers the universal softmax inflection
84 baseline, and the fitted slope of 0.28 is roughly half the $\sqrt{d}$ coefficient predicted for random patterns, reflecting
85 the structured similarity spectrum of real protein sequences arising from conserved positions, correlated
86 substitutions, and phylogenetic signal. Adding further predictors (mean column entropy, effective number
87 of sequences, spectral concentration) produced negligible improvement ($\Delta R^2 < 0.001$), indicating that PCA
88 dimensionality captures most of the relevant variation. Refitting on the eight Pfam families alone yielded a
89 nearly identical relationship ($R^2=0.95$), confirming that the result is not driven by overrepresentation of
90 the WW family. This regression provides a practical shortcut: given a new family, one can estimate β^*
91 from PCA alone and set the generation temperature without a sweep.

92 We compared stochastic attention against three established learned baselines, profile HMM emit, EvoDiff
93 (MSA-conditioned diffusion), and the MSA Transformer (Gibbs masked language model sampling), across
94 all eight families (Fig. 3). Reading the three panels jointly is essential: no single metric tells the full story,
95 and a method that scores well on novelty while failing on compositional fidelity or sequence identity has
96 simply generated sequences that no longer belong to the family.

97 Profile HMM emit and the MSA Transformer achieved high novelty scores, but this apparent advantage
98 is illusory: their nearest sequence identity to stored patterns fell to 11–59% in several families, well below
99 the typical within-family range of 30–70% (shaded band, Fig. 3C). Sequences below this range have drifted
100 outside the family and are unlikely to retain its function; sequences above it are near-copies of stored
101 patterns with insufficient novelty. The high novelty of HMM emit and the MSA Transformer reflects
102 compositional drift, not meaningful exploration of family sequence space. EvoDiff balanced the metrics
103 more effectively in shorter families but required pretraining on millions of UniRef50 MSAs, scaled poorly
104 with sequence length (~ 14 hours for 150 Pkinase sequences on CPU, though GPU inference would reduce
105 this substantially, versus seconds for SA on the same hardware; SI Appendix, Table S10), and suffered a
106 sharp degradation in compositional fidelity on the longest family (Pkinase KL=0.20, an order of magnitude
107 worse than SA).

108 Stochastic attention is the only method that simultaneously achieves all three desiderata: KL divergence
109 below 0.05 in seven of eight families (a KL of 0.05 over 20 amino acids corresponds to an average absolute
110 frequency error of less than 0.5 percentage points per residue type, well within finite-sample variation),
111 novelty of 0.40–0.65 indicating genuine departure from stored patterns, and sequence identity of 51–66%
112 firmly within the natural within-family range. This combination (compositionally faithful, genuinely novel,
113 and staying within the family) is what distinguishes SA from every baseline we tested, using only the
114 seed alignment and no pretraining, backpropagation, or GPU resources. A systematic comparison with a
115 Potts model fitted via pseudo-likelihood maximization across all eight families (SI Appendix, Table S6)
116 revealed comparable compositional fidelity and broadly similar output quality in six of eight families, with
117 SA achieving higher novelty in seven of eight. The Potts model showed clear failure modes in the most
118 data-starved families (Pkinase: low novelty; Defensin_beta: poor compositional fidelity), consistent with its
119 extreme over-parameterization, while SA remained stable. These results suggest that the Hopfield energy
120 implicitly captures the correlational structure that pairwise coevolutionary models fit explicitly, but without
121 parameter estimation or the associated risk of overfitting.

122 We assessed structural plausibility by predicting folded structures for 50 sequences per method per family

123 using ESMFold (Fig. 4), comparing SA generation to stored natural sequences via Wilcoxon rank-sum tests
 124 with Cohen's d effect sizes. SA-generated sequences achieved mean predicted local distance difference test
 125 (pLDDT) scores statistically indistinguishable from stored sequences in six of eight families ($p > 0.05$), with
 126 significantly higher pLDDT in Kunitz (90.8 vs. 89.3, $p=0.001$, $d=0.67$), and across all families at least 86%
 127 of SA-generated sequences exceeded the pLDDT=70 confidence threshold. Template modeling (TM) scores
 128 showed a consistent pattern: in six of eight families, SA generation achieved significantly higher TM-scores
 129 than stored sequences ($p < 0.05$; $d > 1.0$ in RRM, WW, PDZ, and Pkinase), likely reflecting the Boltzmann
 130 ensemble's tendency to weight the central tendency of the family distribution, producing sequences closer
 131 to a consensus fold than individual natural sequences, which carry lineage-specific structural deviations. As
 132 an independent cross-validation, we repeated the structure prediction for all eight families using AlphaFold2
 133 (SI Appendix, Figs. S3–S4). The AF2 results corroborated the ESMFold findings while resolving the single
 134 negative result: for SH3, the only family where ESMFold had indicated a pLDDT deficit, AlphaFold2
 135 showed no pLDDT difference ($p=0.99$) and significantly higher TM-scores for SA generation (0.736 vs.
 136 0.721, $p < 0.001$). The concordance between two independent structure predictors across all eight families
 137 strengthens the computational evidence that SA-generated sequences are predicted to adopt the correct
 138 fold (Fig. 6).

139 The preceding analyses established that SA preserves single-site statistics and structural plausibility;
 140 we next asked whether this fidelity extends to pairwise residue covariation, which reflects structural
 141 contacts and functional couplings that single-site statistics cannot capture. We computed pairwise mutual
 142 information (MI) matrices for the stored and SA-generated alignments in each family and found that
 143 the Pearson correlation between stored and SA-generated MI vectors ranged from $r=0.53$ to 0.92 across
 144 seven of eight families, with Pkinase ($r=0.05$) a consistent outlier for all methods owing to its extreme
 145 data-to-position ratio ($K=37$, $L=262$). SA consistently exceeded profile HMMs ($r=0.07$ –0.83; SI Appendix,
 146 Table S3), and the general ordering $\text{HMM} < \text{SA} < \text{Potts} \leq \text{EvoDiff}$ held across most families, consistent with
 147 the information hierarchy of each approach: position-independent frequencies, implicit all-order Hopfield
 148 couplings, explicitly fitted pairwise couplings, and deep coevolutionary priors from large-scale pretraining,
 149 respectively. This ordering confirms that the Hopfield energy captures correlational structure beyond
 150 what position-independent models encode, without requiring explicit pairwise parameter estimation. As a
 151 complementary, sequence-level assessment of biological plausibility, we scored all SA-generated and stored
 152 sequences using ESM2-650M, a protein language model pretrained on millions of UniRef50 sequences that
 153 is entirely independent of the stochastic attention framework. SA-generated sequences were statistically
 154 indistinguishable from natural family members in four of eight families, with SA achieving significantly
 155 better pseudo-perplexity in the Kunitz domain (4.47 vs. 4.98, $p < 0.001$). Across all eight families, the
 156 mean pseudo-perplexity was 5.71 for SA-generated sequences versus 5.58 for stored sequences, confirming
 157 that an independent model trained on billions of natural sequences considers SA-generated sequences
 158 comparably protein-like to natural family members. We further cross-referenced SA-generated substitutions
 159 against published deep mutational scanning datasets for the PDZ and SH3 domains, finding that 91.7%
 160 and 90.4% of matched substitutions corresponded to experimentally tolerated mutations, representing
 161 $1.18\times$ and $1.34\times$ enrichment over the overall dataset tolerance rates of 77.5% and 67.5%, respectively
 162 (binomial $p < 10^{-135}$ and $p < 10^{-156}$; SI Appendix, Tables S4 and S5). These results demonstrate that the
 163 Boltzmann sampling procedure preferentially generates substitutions at positions and to residues that are
 164 experimentally functional, without any explicit fitness objective.

165 Finally, the beta defensin family (PF00711) provided an opportunity to test whether SA generation
 166 preserves not only sequence-level statistics but also the biophysical properties that underpin antimicrobial
 167 function (Fig. 5). SA-generated sequences preserved the six-cysteine disulfide scaffold (6.01 ± 0.08 mean
 168 cysteines, compared to 6.02 ± 0.15 for natural sequences), maintained a cationic charge profile ($+6.1 \pm 1.6$)
 169 with the narrowest charge distribution of any method, and occupied the same region of charge-hydrophobicity
 170 space as natural defensins. By a minimal antimicrobial plausibility filter (net charge $\geq +2$, hydrophobic
 171 fraction in $[0.3, 0.7]$, ≥ 4 cysteines), 51% of SA-generated sequences passed, compared to 42% of stored

172 sequences, 25% of HMM-emitted sequences, and 16% of MSA Transformer sequences. This implicit
173 enforcement of biophysical constraints arises because the energy landscape concentrates probability near the
174 stored patterns, which already satisfy these constraints by virtue of being functional proteins; no explicit
175 objective for charge, hydrophobicity, or disulfide topology is required.

176 Discussion

177 We set out to ask what can be generated from a small protein family alignment with no training. The answer:
178 sequences that are novel, compositionally faithful, and predicted by independent structure predictors to
179 adopt the correct three-dimensional fold, produced in seconds on a laptop from nothing but the alignment
180 itself. Across eight Pfam families spanning a ten-fold range of family sizes ($K=37-420$) and a seven-fold
181 range of sequence lengths ($L=23-262$), stochastic attention produced output quality comparable to methods
182 pretrained on millions of sequences, while requiring zero external data, zero training, and zero GPU resources.
183 This combination of properties was not achieved by any single learned baseline we tested. Profile HMMs
184 and the MSA Transformer generate sequences with sequence identity well below the natural within-family
185 range (30–70%), indicating that high novelty scores for these methods reflect compositional drift rather
186 than meaningful exploration. EvoDiff requires millions of pretraining sequences and hours of compute,
187 and collapses on the longest family. No baseline simultaneously occupied the within-family identity range,
188 maintained low composition divergence, and produced genuine novelty; stochastic attention is the only
189 method that does. The fact that a training-free method, one that treats the seed alignment as both the
190 model and the data, produces output comparable to systems built on orders of magnitude more information
191 suggests that the statistical structure of protein families is rich enough to support generation without learned
192 representations, provided the generative framework respects the geometry of the family's sequence space.
193 Two control experiments confirmed that SA captures more than marginal statistics: a consensus-with-noise
194 baseline that adds isotropic Gaussian perturbations in PCA space produced a qualitatively different output
195 profile (SI Appendix, Table S11), and SA run on column-permuted alignments (which preserve per-position
196 frequencies but destroy inter-position correlations) yielded consistently different novelty and sequence
197 identity patterns (SI Appendix, Table S12).

198 The structure validation results were noteworthy. SA-generated sequences achieved significantly higher
199 TM-scores to the canonical family fold in six of eight families by ESMFold, and this finding was corroborated
200 by AlphaFold2 across all eight families. The pLDDT confidence scores were statistically indistinguishable
201 from stored sequences in six of eight families, confirming that the sampler does not sacrifice folding
202 confidence for diversity, and the single exception (SH3) was resolved by AlphaFold2 cross-validation. Two
203 explanations for the elevated TM-scores are possible: structure predictors, trained on evolutionary data,
204 may assign higher confidence to consensus-proximal sequences, or the consensus-proximal sequences may
205 genuinely adopt a fold closer to the canonical family structure. Since the Boltzmann ensemble weights
206 the central tendency of the family distribution, the generated sequences are inherently closer to consensus
207 than individual natural sequences, making the predictor-bias explanation plausible. We note that all
208 computational validators used in this study (structure predictors, ESM2-650M, and DMS-derived fitness
209 scores) are trained on evolutionary data, so high scores for family-like sequences are expected; experimental
210 structure determination of selected generated sequences would provide definitive confirmation.

211 The scaling study on the WW domain showed that stochastic attention is not fragile: generation quality
212 was stable down to $K=20$ sequences, with KL divergence below 0.02 and novelty above 0.47 even at this
213 extreme, and replication across three additional families confirmed generality. This robustness arises from a
214 structural property of the method: the energy landscape is defined entirely by the stored patterns, with
215 no parameters that could overfit to a small training set. There is no encoder to memorize, no decoder to
216 collapse, and no denoising network to undertrain. The entropy inflection criterion automatically adapts the
217 operating temperature to the geometry of the memory matrix at each scale, and the critical temperature β^*
218 itself can be predicted from PCA dimensionality alone ($R^2=0.97$), eliminating the need for a temperature
219 sweep and enabling fully automatic operation from a seed alignment. In contrast, learned generative models

face a fundamental parameter-to-data ratio problem when $K < 100$: VAEs and diffusion models have thousands to millions of parameters that cannot be meaningfully constrained by a few dozen sequences. This regime encompasses most Pfam families and the long tail of protein space where generative tools are most needed.

Looking beyond proteins, the connection between attention and Boltzmann sampling suggests a broader principle: any set of examples, regardless of domain, defines an energy landscape via the modern Hopfield energy, and Langevin dynamics on this landscape provides a principled generative model with no training. The requirements are minimal: the examples must be representable as vectors in a common space, and the practitioner must choose an inverse temperature that balances fidelity and diversity. This perspective may extend naturally to other biological sequence families where data is scarce but examples are structured, such as antibody repertoires where individual clonotypes may have only a handful of somatic variants, RNA families where Rfam seed alignments are typically smaller than their Pfam counterparts, or small-molecule libraries where chemical series may contain only tens of analogs. The analytic score function also provides a natural interface with the convergence theory of Langevin dynamics (22, 23), enabling formal guarantees on sample quality that are absent from most neural generative models and opening a path toward certifiable generation in safety-critical applications such as therapeutic protein design.

Several limitations should be acknowledged. The most significant is that stochastic attention operates on a fixed alignment representation: it requires a pre-computed MSA and a fixed alignment length L , and cannot generate insertions, deletions, or variable-length sequences. This constraint is inherited from the one-hot-plus-PCA encoding, which maps each sequence to a vector of fixed dimension $20L$; any change in L would alter the representation space itself. In practice, this means the method is best suited to compact, well-aligned domains rather than multi-domain proteins or families with frequent large insertions. Extending stochastic attention to variable-length generation would likely require replacing PCA with a length-agnostic embedding while preserving the closed-form score function, a nontrivial but tractable direction for future work. The PCA projection itself discards information that may be relevant for capturing higher-order correlations, though a sensitivity analysis confirmed robustness across the 75%–99% retained variance range (SI Appendix, Table S13). The eight families studied here are all well-structured, globular domains with abundant structural and functional annotation; this choice was necessary for validation but introduced a selection bias toward precisely the class of proteins that structure predictors handle best. Performance on intrinsically disordered regions, transmembrane domains, multi-domain proteins, or families with very low sequence identity remains untested, and the method’s reliance on PCA may be particularly limiting for disordered families where sequence variation is less structured. Finally, our multi-chain protocol does not guarantee coverage of all modes in the Boltzmann distribution, particularly for very large or highly diverse families, and wet-lab characterization of selected generated sequences would be needed to establish the method’s utility for protein engineering applications.

Materials and Methods

Stochastic Attention. Let $\mathbf{X} = [\mathbf{m}_1, \dots, \mathbf{m}_K] \in \mathbb{R}^{d \times K}$ be the memory matrix whose columns are vectorized protein sequences from a family alignment of K members, each represented in d dimensions after one-hot encoding and PCA projection. The modern Hopfield energy (19) is:

$$E(\boldsymbol{\xi}) = \frac{1}{2} \|\boldsymbol{\xi}\|^2 - \frac{1}{\beta} \log \sum_{k=1}^K \exp(\beta \mathbf{m}_k^\top \boldsymbol{\xi}), \quad [1]$$

where $\beta > 0$ is an inverse temperature. The Boltzmann distribution $p_\beta(\boldsymbol{\xi}) \propto \exp(-\beta E(\boldsymbol{\xi}))$ has score function $\nabla \log p_\beta(\boldsymbol{\xi}) = \beta(\mathbf{T}(\boldsymbol{\xi}) - \boldsymbol{\xi})$, where $\mathbf{T}(\boldsymbol{\xi}) := \mathbf{X} \text{softmax}(\beta \mathbf{X}^\top \boldsymbol{\xi})$ is the softmax attention retrieval map. This score is *exact* and computed by a single attention operation; no score network or training is required. Applying the unadjusted Langevin algorithm (ULA) (24) to the potential $U = \beta E$ yields the stochastic attention update:

$$\boldsymbol{\xi}_{t+1} = (1 - \alpha) \boldsymbol{\xi}_t + \alpha \mathbf{X} \text{softmax}(\beta \mathbf{X}^\top \boldsymbol{\xi}_t) + \sqrt{\frac{2\alpha}{\beta}} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{I}). \quad [2]$$

Each step performs contraction toward the origin, a softmax-weighted pull toward stored memories, and isotropic Gaussian noise, at a per-step cost of $\mathcal{O}(dK)$.

Temperature Selection and Sequence Decoding. The critical inverse temperature β^* is identified via the entropy inflection criterion: we evaluate the Boltzmann entropy $H(\beta) = -\sum_k w_k(\beta) \log w_k(\beta)$, where $w_k(\beta) = \text{softmax}(\beta \mathbf{X}^\top \boldsymbol{\xi})_k$ are the attention weights at a representative stored pattern $\boldsymbol{\xi} = \mathbf{m}_1$, and locate the inflection point $\beta^* = \arg \max_\beta |d^2 H/d\beta^2|$ via a grid search over $\beta \in [0.1, 20]$. Below β^* the sampler visits all patterns equiprobably; above β^* pattern-specific basins emerge. Generation uses $\beta_{\text{gen}} \approx 2\beta^*$, just above the inflection, to balance exploration and fidelity.

To decode a PCA-space sample $\boldsymbol{\xi}$ to an amino acid sequence, the PCA projection is inverted to recover a one-hot-like vector in the original $20L$ -dimensional space, and the amino acid at each position is assigned by argmax over the 20 channels. The resulting length- L sequence is the decoded output; positions where the PCA reconstruction assigns negligible weight to all amino acids default to the most frequent residue in the stored alignment at that position.

Experimental Design. We selected eight Pfam families spanning a range of sizes, sequence lengths, and conservation levels (Table 1): RRM (PF00076, $K=68$, $L=71$), SH3 (PF00018, $K=55$, $L=48$), WW (PF00397, $K=420$, $L=31$), Kunitz (PF00014, $K=99$, $L=53$), zf-C2H2 (PF00096, $K=151$, $L=23$), PDZ (PF00595, $K=44$, $L=83$), Pkinase (PF00069, $K=37$, $L=262$), and Defensin_beta (PF00711, $K=45$, $L=36$). Seed alignments were downloaded from InterPro (18) and cleaned by removing columns with $> 50\%$ gaps and sequences with $> 30\%$ gaps. PCA retained 95% of variance, followed by unit-norm projection to form the memory matrix. For each family, we ran 30 independent ULA chains of $T=5,000$ iterations ($\alpha=0.01$), discarding 2,000 burn-in iterations and thinning every 100 steps to retain 5 samples per chain ($S=150$ sequences total). The inverse temperature was set automatically via the entropy inflection criterion, with a generation regime at $\beta_{\text{gen}} \approx 2\beta^*$ and a retrieval regime at $\beta_{\text{ret}} \approx 20\beta^*$. The thinning interval of 100 steps is well-matched to the integrated autocorrelation time of the Hopfield energy ($\tau_{\text{int}} = 76\text{--}122$ steps across families), yielding approximately 35 effective independent samples per chain and $\approx 1,050$ effective draws total; the burn-in of 2,000 steps provides a $3.2\text{--}6.6\times$ margin over the empirical chain convergence point (SI Appendix, Table S10).

To contextualize SA against methods that require pretraining on large external corpora, we compared four established baselines: (i) profile HMM emit via HMMER3 (5); (ii) EvoDiff MSA-conditioned diffusion (15), pretrained on UniRef50 MSAs; (iii) MSA Transformer (14) via iterative Gibbs-like masked language model sampling; and (iv) a Potts model fitted via pseudo-likelihood maximization (plmDCA). For each baseline, the same 150-sequence budget and evaluation metrics (KL divergence, novelty, sequence identity) were applied. Structural plausibility was assessed by predicting folded structures for 50 sequences per method per family using ESMFold (25) and AlphaFold2 (1) (via ColabFold (26)), comparing each predicted structure to an experimentally determined reference using TM-align (27). All statistical comparisons used Wilcoxon rank-sum tests with Benjamini-Hochberg false discovery rate (FDR) correction at $q=0.05$.

Acknowledgments. The idea for stochastic attention arose while developing a lecture on energy-based methods for eCornell. We thank the Pfam/InterPro consortium for maintaining publicly available seed alignments. Structure predictions were performed using ESMFold and AlphaFold2 (via ColabFold). Computations were performed on resources at Cornell University.

Data, Materials, and Software Availability

All code, data, seed alignments, generated sequences, and analysis scripts are publicly available at <https://github.com/varnerlab/SA-Protein-Modeling-Study> (21). Pfam seed alignments were obtained from InterPro (18). No new experimental data were generated for this study; all validation used publicly available computational tools (ESMFold, AlphaFold2 via ColabFold, ESM2-650M) and published deep mutational scanning datasets.

1. Jumper J, et al. (2021) Highly accurate protein structure prediction with AlphaFold. *Nature* 596:583–589.
2. Baek M, et al. (2021) Accurate prediction of protein structures and interactions using a three-track neural network. *Science* 373(6557):871–876.
3. Kuhlman B, Bradley P (2019) Advances in protein structure prediction and design. *Nature Reviews Molecular Cell Biology* 20:681–697.
4. Huang PS, Boyken SE, Baker D (2016) The coming of age of de novo protein design. *Nature* 537(7620):320–327.
5. Potter SC, et al. (2018) HMMER web server: 2018 update. *Nucleic Acids Research* 46(W1):W200–W204.
6. Krogh A, Brown M, Mian IS, Sjölander K, Haussler D (1994) Hidden Markov models in computational biology: Applications to protein modeling. *Journal of Molecular Biology* 235(5):1501–1531.
7. Morcos F, et al. (2011) Direct-coupling analysis of residue coevolution captures native contacts across many protein families. *Proceedings of the National Academy of Sciences* 108(49):E1293–E1301.
8. Ekeberg M, Lökvist C, Lan Y, Weigt M, Aurell E (2013) Improved contact prediction in proteins: Using pseudolikelihoods to infer Potts models. *Physical Review E* 87(1):012707.
9. Trinquier J, Uguzzoni G, Pagnani A, Zamponi F, Weigt M (2021) Efficient generative modeling of protein sequences using simple autoregressive models. *Nature Communications* 12:5800.
10. Kingma DP, Welling M (2014) Auto-encoding variational Bayes in *International Conference on Learning Representations (ICLR)*.

- 318 11. Riesselman AJ, Ingraham JB, Marks DS (2018) Deep generative models of genetic variation capture the effects of mutations. *Nature Methods* 15:816–822.
- 319 12. Ferruz N, Schmidt S, Höcker B (2022) ProtGPT2 is a deep unsupervised language model for protein design. *Nature Communications* 13:4348.
- 320 13. Madani A, et al. (2023) Large language models generate functional protein sequences across diverse families. *Nature Biotechnology* 41:1099–1106.
- 321 14. Rao R, et al. (2021) MSA transformer in *Proceedings of the 38th International Conference on Machine Learning (ICML)*. pp. 8844–8856.
- 322 15. Alamdari S, et al. (2023) Protein generation with evolutionary diffusion: Sequence is all you need. *bioRxiv*.
- 323 16. Dauparas J, et al. (2022) Robust deep learning–based protein sequence design using ProteinMPNN. *Science* 378(6615):49–56.
- 324 17. Watson JL, et al. (2023) De novo design of protein structure and function with RFdiffusion. *Nature* 620(7976):1089–1100.
- 325 18. Mistry J, et al. (2021) Pfam: The protein families database in 2021. *Nucleic Acids Research* 49(D1):D412–D419.
- 326 19. Ramsauer H, et al. (2021) Hopfield networks is all you need in *International Conference on Learning Representations (ICLR)*.
- 327 20. Krotov D, Hopfield JJ (2016) Dense associative memory for pattern recognition in *Advances in Neural Information Processing Systems*. Vol. 29.
- 328 21. Alswaidan A, Varner JD (2026) Stochastic attention via Langevin dynamics on the modern Hopfield energy. *arXiv preprint arXiv:2603.06875*.
- 329 22. Durmus A, Moulines É (2017) Nonasymptotic convergence analysis for the unadjusted Langevin algorithm. *The Annals of Applied Probability* 27(3):1551–1587.
- 330 23. Dalalyan AS (2017) Theoretical guarantees for approximate sampling from smooth and log-concave densities. *Journal of the Royal Statistical Society: Series B* 79(3):651–676.
- 331 24. Roberts GO, Tweedie RL (1996) Exponential convergence of Langevin distributions and their discrete approximations. *Bernoulli* 2(4):341–363.
- 332 25. Lin Z, et al. (2023) Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science* 379(6637):1123–1130.
- 333 26. Mirdita M, et al. (2022) ColabFold: Making protein folding accessible to all. *Nature Methods* 19:679–682.
- 334 27. Zhang Y, Skolnick J (2005) TM-align: a protein structure alignment algorithm based on the TM-score. *Nucleic Acids Research* 33(7):2302–2309.

Significance Statement

Most protein families have fewer than 100 known members, too few to train a deep generative model. We show that a single attention operation over a family alignment defines an energy landscape whose Boltzmann distribution can be sampled via Langevin dynamics, with no training, no GPU, and no external data. The resulting method, stochastic attention, generates novel protein sequences that preserve family-level amino acid statistics and are predicted by independent structure predictors to adopt the correct fold across eight diverse families. A simple linear relationship between representation dimensionality and the critical sampling temperature enables fully automatic operation from a seed alignment, opening training-free sequence generation to the majority of protein families too small for deep learning.

J.D.V. designed research, performed research, analyzed data, and wrote the paper.

The author declares no competing interest.

¹ To whom correspondence should be addressed. E-mail: jdv27@cornell.edu. ORCID: 0000-0002-2558-7026.

Table 1. Properties of the eight Pfam families used in this study. K : number of seed sequences after cleaning. L : alignment length. d : PCA dimension (95% variance). $\bar{H}_{\text{col}}$: mean per-column Shannon entropy (nats). K_{eff} : effective number of sequences at 80% identity. β^* : empirical entropy inflection point.

Family	Pfam ID	K	L	d	$\bar{H}_{\text{col}}$	K_{eff}	β^*	$\beta^* / \sqrt{d}$
RRM	PF00076	68	71	59	1.92	68	3.85	0.50
SH3	PF00018	55	48	46	1.68	55	3.85	0.57
WW	PF00397	420	31	186	1.74	419	5.45	0.40
Kunitz	PF00014	99	53	80	1.61	99	3.85	0.43
zf-C2H2	PF00096	151	23	106	1.91	151	4.58	0.44
PDZ	PF00595	44	83	37	1.88	43	3.23	0.53
Pkinase	PF00069	37	262	34	1.72	37	3.23	0.55
Defensin_beta	PF00711	45	36	37	1.65	45	3.23	0.53




Fig. 1. Cross-family comparison of SA generation, SA retrieval, and bootstrap across eight Pfam families. **(A)** KL divergence of amino acid composition (lower is better); dashed line at $KL=0.05$ marks the practical threshold below which deviations are indistinguishable from finite-sample variation. SA generation achieves the lowest or near-lowest KL in every family. **(B)** Novelty in PCA space ($1 - \max_k \cos(\hat{\xi}, \mathbf{m}_k)$; higher is better). SA generation produces substantially novel sequences, while bootstrap samples, which are exact copies of stored patterns, have zero PCA-space novelty by construction. **(C)** Nearest sequence identity to a stored pattern; the green shaded band (30–70%) marks the typical within-family identity range; sequences below this band have drifted outside the family, while sequences above it are near-copies of stored patterns with insufficient novelty. SA generation falls within this range across all eight families. Error bars show ± 1 SE (30 chains $\times$ 5 samples).




Fig. 3. Head-to-head comparison of SA generation against three learned baselines (profile HMM emit, EvoDiff, MSA Transformer) across all eight Pfam families. **(A)** KL divergence of amino acid composition (pooled over all 150 generated sequences, with bootstrap SE; $n_{\text{boot}}=1000$); dashed line at $\text{KL}=0.05$ marks the threshold below which deviations are indistinguishable from finite-sample variation. **(B)** Novelty in PCA space. **(C)** Nearest sequence identity to a stored pattern; the green shaded band (30–70%) marks the typical within-family identity range; sequences below this band have drifted outside the family, while sequences above it are near-copies of stored patterns with insufficient novelty. Panels (B) and (C) report per-chain means ± 1 SE (30 chains $\times$ 5 samples). Significance brackets show Wilcoxon rank-sum tests comparing SA generation to each baseline (HMM emit, EvoDiff, MSA Transformer); only significant comparisons are shown ($***p < 0.001$, $**p < 0.01$, $*p < 0.05$). Simple reference baselines (Gaussian perturbation, convex combination, stored-pattern retrieval) are omitted from this figure because their design precludes meaningful comparison; their metrics are reported in SI Appendix, Table S1 for completeness.




Fig. 4. Structure validation of SA-generated sequences using ESMFold. **(A)** Mean pLDDT (folding confidence) across all eight Pfam families for SA generation, SA retrieval, bootstrap, and stored (natural) sequences. The dashed line indicates the pLDDT=70 confidence threshold. **(B)** TM-score to a representative experimentally determined structure for each family. The dashed line indicates TM=0.5 (same-fold threshold). Significance brackets show Wilcoxon rank-sum tests comparing SA generation to stored sequences: *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, n.s. = not significant. Error bars show ± 1 SE ($n=50$ sequences per condition).




Fig. 5. Biophysical properties of beta defensin sequences (PF00711). **(A)** Net charge at pH 7 vs. hydrophobic residue fraction. SA-generated sequences (blue) overlap the stored natural sequences (gray), while HMM emit (green) and MSA Transformer (purple) drift to lower charge. Dashed lines indicate the antimicrobial plausibility thresholds (charge $\geq +2$, hydrophobic fraction 0.3–0.7). **(B)** Cysteine count distribution. SA generation preserves the six-cysteine disulfide scaffold (red dashed line) with the same fidelity as the stored alignment. **(C)** Net charge distribution. SA generation produces a charge profile closely matching stored sequences. **(D)** Fraction of sequences passing a minimal biophysical plausibility filter (charge $\geq +2$, hydrophobic fraction $\in [0.3, 0.7]$, ≥ 4 cysteines).




Fig. 6. AlphaFold2 predicted structures for representative SA-generated sequences (colored by per-residue pLDDT confidence) superimposed on the experimentally determined reference structure for each family (gray, semi-transparent). For each family, the SA-generated sequence with the highest mean pLDDT was selected and aligned to the reference using PyMOL. The close predicted structural agreement (sub-angstrom root-mean-square deviation (RMSD)) and uniformly high pLDDT scores are consistent with SA-generated sequences encoding three-dimensional folds that recapitulate the canonical architecture of their target family, despite 40–65% sequence novelty.